

Dynamic Resource Exchange Performance in Cyclus

Matthew J. Gidden, Paul P. H. Wilson

University of Wisconsin, Madison WI
gidden@wisc.edu

INTRODUCTION

Nuclear fuel cycle simulation is a field which seeks to model the facilities and material flows required to produce nuclear power. Simulations normally model a time span of decades, or even centuries. Myriad decisions exist within a given fuel cycle simulation, such as the deployment timing of facility types and determining how material transfers should be executed at a given time step. Accordingly, trade-offs exist between the features provided by a simulator and the performance of the simulator.

The CYCLUS simulator [1] was designed to more easily model a variety of fuel cycles. It uses an agent-based modeling paradigm, encapsulating the difficulty of designing new agent archetypes. Archetypes define parameterized agent logic and behavior and can therefore be reused within and between simulations. The core agent-interaction model in CYCLUS is the Dynamic Resource Exchange (DRE) [2, 3]. The DRE, recomputed at each time step, polls the supply and demand of commodities in the simulation and then determines the trades to be executed between agents. Highly dynamic, easily adjustable fuel cycles can be modeled by coupling the DRE concept with the CYCLUS Region-Institution-Facility hierarchy and archetype-prototype-agent models, addressing many of the issues developers and users have found with previous models. However, CYCLUS must both be featureful and performant. This work provides a first-look at how the DRE supply-demand model scales with problem size by generating and solving a large number of exchanges. Different solvers are analyzed including a CYCLUS-aware greedy heuristic in addition to COIN-OR's LP and MILP solvers [4, 5].

METHODOLOGY

Instances of resource exchanges are required to analyze the effects and performance of the DRE and its solvers. In the absence of large Cyclus simulations with interesting facility and relationship models, instances must be generated given some set of rules and parameters. Two distinct *species* of exchanges are generated, those related to the *front end* of the nuclear fuel cycle and those related to the *back end* of the nuclear fuel cycle. For each species, commodities, reactors, and supporting facilities must be generated in addition to exchange-related parameters, such as the trade preferences.

Three types of fuel cycles are generated: a once-through fuel cycle, labeled OT; a plutonium-recycle fuel cycle, labeled MOX; and a plutonium and thorium-recycle fuel cycle, labeled MOX-ThOX. As fuel cycles increase in complexity, the number of commodities that exist increases, as shown in Table I. The commodities are referred to by abbreviation: Uranium Oxide (UOX), Mixed Plutonium Oxide for Thermal Reactors (TMOX), Mixed Plutonium Oxide for Fast Reactors (FMOX),

TABLE I. A mapping between fuel cycles to the commodities that exist in each one.

| Fuel Cycle | Commodities |
|------------|------------------------------|
| OT | UOX |
| MOX | UOX TMOX FMOX |
| ThOX | UOX TMOX FMOX FThOX |

TABLE II. A mapping between reactor types and the commodities allowed to fuel each reactor type.

| Reactor Types | Fuel Commodities |
|---------------|------------------------------|
| Thermal | UOX TMOX FMOX |
| FMOX | UOX TMOX FMOX FThOX |
| FThOX | UOX TMOX FMOX FThOX |

Thorium Oxide for Fast Reactors (FThOX).

All reactors are modeled as either thermal or fast reactors. It is necessary to estimate the amount of fuel exchanged by reactors each time step. Accordingly, thermal reactors are simplified models of AP-1000 reactors, and fast reactors are simplified models of BN-600 reactors. Reactors operate in a batch mode, where each batch is approximately one quarter of the reactor core, an assumption used by other analyses [6]. Each reactor in the exchange is either offering or requesting a batch of fuel. Reactors may be fueled by different fuel types, i.e., fuel commodities. A mapping of reactors to acceptable commodities is provided in Table II. Note that there is still a preference distribution associated with each reactor-commodity pair as well as constraint coefficient effects; therefore, while a reactor *can* consume a fuel commodity, it may do so only in the absence of any other commodity availability. It is assumed that a reactor would rather utilize an undesirable fuel type than realize the lost revenue from not producing power.

In a front-end exchange, fuel suppliers exchange material with reactors. In a back-end exchange, reprocessing and storage facilities exchange material with reactors. In either

in this simulation
this is not a constraint of Cyclus

It is not clear whether this section is describing the constraints of a simulation to be run or the constraints of Cyclus.

I think you are describing the former. Perhaps a simple clarifying sentence would help

case, facilities that are not reactors are referred to as *support* facilities, as they support the reactors which generate power. In the front-end exchange model, a single type of supporting facility exists for each fuel type, e.g., an enrichment facility for UOX, a thermal reprocessing plant for TMOX, etc. In the back-end exchange model, a similar supporting facility assignment is used, except a storage facility type, capable of storing any commodity is added.

For reactors and support facilities in both exchange species, exchange constructs, such as preferences and constraint coefficients, must be generated. Constraint coefficients are a function of resource quality, for which a proxy must be assigned. Accordingly, each reactor is modeled using an enrichment-level-range-to-commodity mapping. For example, LWRs requesting UOX have an enrichment range of [3.5, 5.5]. The enrichment value for each reactor is chosen from a uniform distribution in order to differentiate reactors. Preferences in the DRE are requester-based, and are assigned per-commodity to requesters in both exchange species. In order to differentiate between preferences, facilities are assigned a location proxy. The location proxy can provide additional weight to preference values.

Exchange generation is parameterized and thus depends on a parameter vector. While each species has specific parameters, both species share some fundamental parameters, namely: f_{fc} , which fuel cycle is being generated; f_{rx} , whether reactors are modeled using a single batch or a collection of assemblies; and f_{loc} , to what degree location is used in determining preference coefficients. f_{fc} may take on values of OT, MOX, or THOX, as described above. f_{rx} may take on values of Batch or Assembly. f_{loc} may take on values of None, Coarse, or Fine.

Given a set of parameters, an exchange instance is generated and persisted in a database. The generated exchange exists in the *exchange layer* of the DRE, as shown in Figure 1. Because the greedy heuristic is knowledgeable of CYCLUS, it uses solution pathway 1 shown in the figure. The COIN solvers must have the CYCLUS exchange structure formulated into a mathematical program, described previously in [2], and thus must take the second solver pathway, through the formulation layer. Figure 1 also notes the time points measured during execution in order to compare solution times between solvers and different exchange instances, i.e., for each solver and exchange instance, execution time is measured as $t_f - t_i$.

In order to explore the large number of possible exchange instances, a sophisticated instance generation and solving framework is needed. Accordingly, a new software package called Cyclopts (Cyclus Optimization Studies) has been developed. Cyclopts provides a general framework for sampling a parameter space, defining problem instances for a given point in parameter space, and solving a problem instance under a variety of conditions. Cyclopts leverages an HTCondor-aware High Throughput Computing (HTC) framework at UW-Madison in order to efficiently run exchange instances.

An experiment consists of a set of resource-exchange graph instances executed with a collection of configured solvers. When a solution is found, the solution (i.e., the flow vector), the time required to reach the solution, and the objective value (i.e., the dot-product of cost and flow vectors)

are recorded. Six execution nodes on UW-Madison Advanced Computing Initiative (ACI) HTCondor system form the homogeneous environment used to conduct the experiments herein described. Each execute node is comprised of an 2.90 GHz eight-core, sixteen-thread, Intel Xeon E5-2690 processor with 128 GB of RAM. Processor hyper-threading was disabled for the duration of the experimental campaign to allow comparisons between solution times.

RESULTS AND ANALYSIS

A scalability campaign was launched to study how the solvers perform as the size of exchange increases. All species-specific parameters for generating an exchange are a function of the number of reactors in the exchange, thus the number of reactors in the system is the natural choice for a scaling parameter. A range of [10, 500] was chosen for this analysis. For each reactor population value, exchanges were generated for all possible combinations of fundamental parameter values (18 total). Each generated exchange instance was solved with all three available solvers. Figure 2 provides an example of output from the scalability campaign.

Both the CLP and greedy solvers were found to scale similarly. As shown in Figure 2, polynomial scaling ($O(n^2)$) in the reactor population is observed. This is driven by the number of constraints in the resulting formulation, rather than the number of variables, for which linear scaling ($O(n)$) was observed. Polynomial scaling of both the greedy and CLP solvers occurs regardless of fundamental parameter or exchange species. The CBC solver, however, performs sporadically, with many exchange instances remaining unconverged at a 3-hour time cap. This is not unexpected, as solving a MILP optimally is an NP-hard problem. Most exchanges consisting of 100 (or less) reactors was found to converge well below the cap.

Every exchange instance depends on both objective value and constraint coefficients, both of which are products of stochastic processes. Accordingly, an experimental campaign was launched to study the effects of stochasticity on the solution times and relative solution values between solvers. An exchange with 65 reactors was chosen to be generated in order to achieve reasonable CBC solve times in most cases. Similar to the scalability campaign, all possible combinations of fundamental parameter values were investigated; 1,000 were generated as a base line case study for each combination. Figure 3 provides an example of output from the stochastic campaign, graphing cumulative solution time values as the number of observations increases.

As can be seen in Figure 3, the CBC solver is sporadic even for relatively small exchange sizes: some instances are solved very quickly, while others can occasionally take a full allotted time before returning. Interestingly, this behavior was seen much more frequently in batch-based exchanges rather than assembly-based exchanges, likely due a more constrained feasible solution space. Stochastic effects were also observed in the solution times for the greedy heuristic and CLP solvers. However, average solution times were found to converge after ~100-200 instances for both species of exchanges and all combinations of fundamental parameters.

this info may be better suited to a table

this may be clearer in a table or graphic

a comment on the practical implications of this scaling behavior would be welcome here.

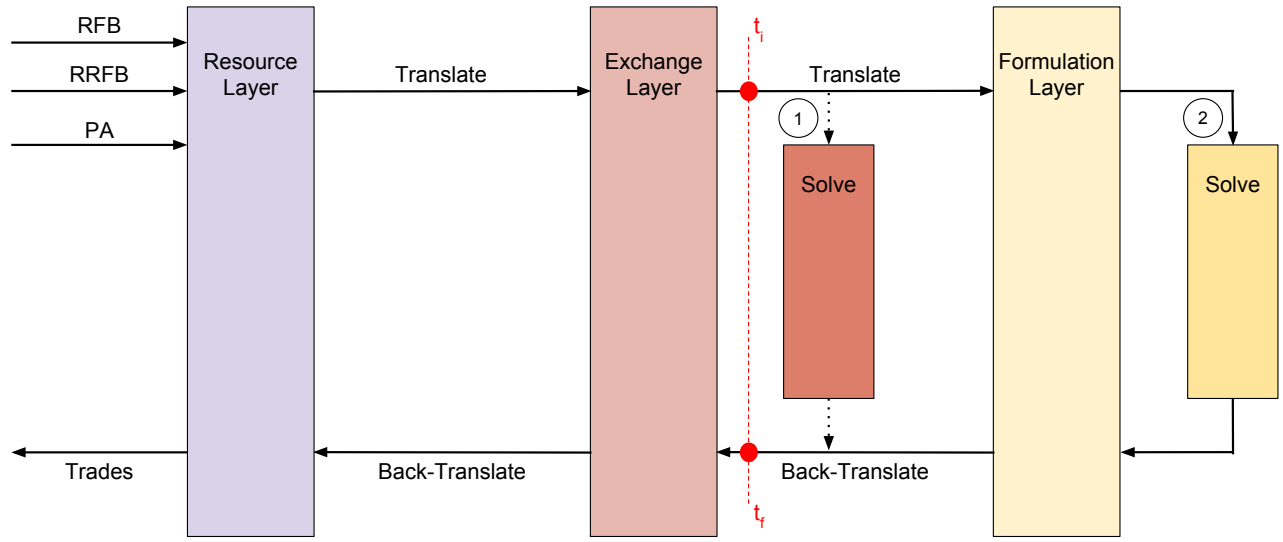


Fig. 1. The control flow of the DRE execution and time points for comparing different solutions. The greedy heuristic takes the pathway labeled 1, while CLP and CBC solvers take the second pathway. The DRE workflow is described more fully in [2, 3].




Fig. 2. Results for the size scoping study for the greedy heuristic modeling a back-end MOX fuel cycle. Rows have increasing f_{tx} values and columns have increasing f_{loc} values.

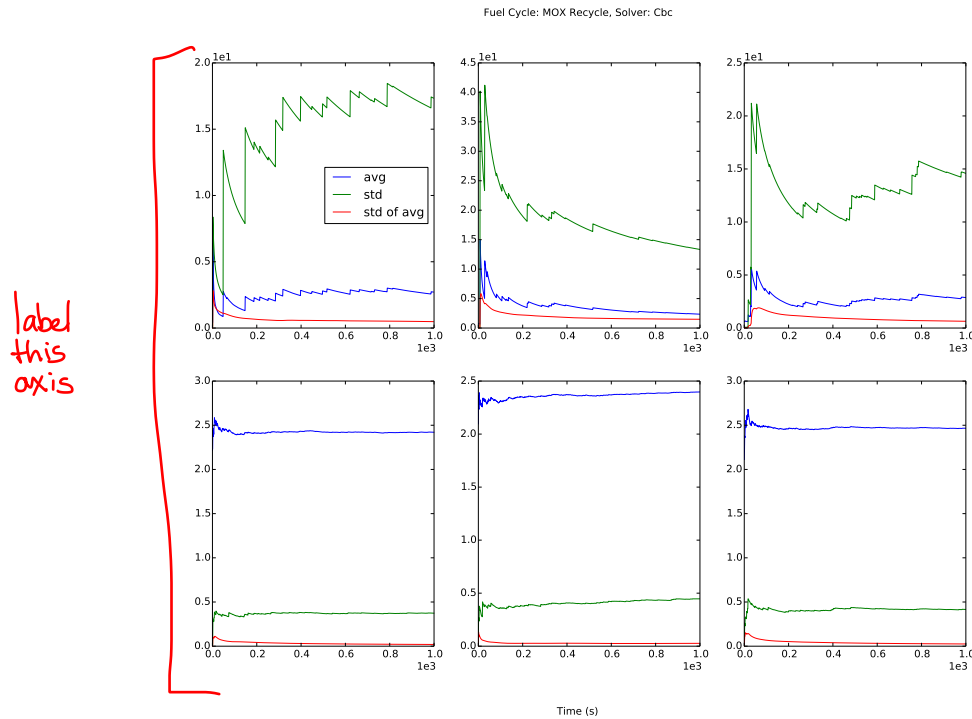


Fig. 3. Timing results for the stochastic study for the CBC Solver modeling a front-end MOX fuel cycle with 65 reactors. Rows have increasing f_{tx} values and columns have increasing f_{loc} values.

CONCLUSIONS

A first-of-a-kind analysis of the dynamic solution of supply and demand in a nuclear fuel cycle is presented. A large number of exchanges are generated and solved with a CYCLUS-aware greedy heuristic, the CLP solver, and the CBC solver. The greedy heuristic is shown to scale well for very large problem sizes, even when modeling up to 500 reactors (on the order of the current world-wide reactor population) trading individual assemblies. The CLP solver also performs well at large problem scales. The CBC solver, while achieving the true optimal solution in many cases, performs sporadically. Such a result is not surprising, especially for highly-constrained optimization problems. Its use, therefore, is suggested for initial exploratory studies before embarking on full-fledged analytic campaigns. Most users will likely find the greedy heuristic to be acceptable, as it provides a feasible solution and enables the powerful features of the DRE.

ACKNOWLEDGMENTS

This research is being performed using funding received from the DOE Office of Nuclear Energy's Nuclear Energy University Programs. The authors thank the NEUP for its generous support.

The authors would like to especially thank Dr. Anthony M. Scopatz, whose advice contributed greatly in the development of Cyclopts.

REFERENCES

1. R. W. CARLSEN, M. GIDDEN, K. HUFF, A. C. OPOTOWSKY, O. RAKHIMOV, A. M. SCOPATZ, Z. WELCH, and P. WILSON, "Cyclus v1.0.0," (06 2014).
2. M. GIDDEN and P. WILSON, "An Agent-Based Framework for Fuel Cycle Simulation with Recycling," in "Proceedings of GLOBAL," Salt Lake City, UT, United States (Sep. 2013).
3. M. GIDDEN, R. CARLSEN, A. OPOTOWSKY, O. RAKHIMOV, A. SCOPATZ, and P. WILSON, "Agent-Based Dynamic Resource Exchange in Cyclus," in "Proceedings of PHYSOR," Kyoto, Japan (Sep 2014).
4. J. FORREST ET AL., "COIN-OR Linear Programming (CLP) Solver," (Dec. 2014).
5. J. FORREST ET AL., "COIN-OR Branch and Cut (CBC) Solver," (Dec. 2014).
6. A. RINEISKI, M. ISHIKAWA, J. JANG, P. MOHANAKRISHNAN, T. NEWTON, G. RIMPAULT, A. STAN-CULESCU, and V. STOGOV, "Reactivity Coefficients in BN-600 Core with Minor Actinides," *Journal of nuclear science and technology*, **48**, 4, 635–645 (2011).